

Database Sharding

- Sharding is **horizontal partitioning** of a database.
- Data is **split** across multiple database instances (shards).
- Each database instance holds a **subset** of the actual data.
- There are multiple ways to split data, such as:
 - **By ID,**
 - **By User,**
 - **By Region, etc.**

Key Features of Database Sharding

- 1) **Horizontal Scaling**
 - Instead of using a single large database,
 - Data is **distributed** across smaller, independent database instances.
- 2) **Partitioned Data**
 - Each database instance holds:
 - **Not the entire dataset,**
 - **but A subset of the total data.**
- 3) **Independent Queries**
 - Instead of running a query on the entire dataset or across all instances,
 - Queries are directed to the **appropriate database instance.**
- 4) **Reduced Load on a Single Database Instance**
 - Prevents overloading a single database, which improves performance.

Purposes of Database Sharding

- When the dataset is **too large** for a single database server to handle.
 - To **improve query performance**
- By reducing the amount of data processed per request.
- To **distribute load** across multiple database servers.
- For **multi-regional applications**
 - Users can retrieve data from the closest database instance.

Benefits of Database Sharding

- 1) **Improved Performance**
 - Smaller subsets of data reduce query execution time.
- 2) **Scalability**
 - New database instances can be added as data and load increase.
- 3) **Fault Tolerance**
 - A failure in one database instance does **not** impact the entire system.

Drawbacks of Database Sharding

- 1) **Increased Code Complexity**
 - Requires:
 - More **complex code**
 - Routing queries to the appropriate database instance.
- 2) **Rebalancing Issues**
 - Some database instances may store **significantly more data** than others.
 - The data needs to be **redistributed** when this happens.
- 3) **Cross-Shard Queries**
 - Joining data **across multiple shards** is difficult.

Example of Database Sharding

- A table has a **numeric ID** column.

- Let's **split** the data as follows:
- Rows with **even IDs** are stored in **Shard 1**.
- Rows with **odd IDs** are stored in **Shard 2**.

```
SELECT * FROM shard_2.user WHERE ID = 1; -- Fetches from Shard 2
SELECT * FROM shard_1.user WHERE ID = 2; -- Fetches from Shard 1
```