

## Performance

- There are two key factors that define the performance of a system:
- 1) **Low Latency**
  - The ability to respond quickly (fast response time).
- 2) **High Availability**
  - The ability to remain accessible (system is always available).

## Low Latency

- **Latency**
  - The time taken by a system to process a request and return a response.
- **Low latency**
  - The system responds quickly (with minimal delay).

## Why is Low Latency Important?

- 1) **Users expect fast interactions**
  - Websites and mobile applications should work quickly.
- 2) **Real-time systems require low latency**
  - Delays in real-time systems can cause serious problems:
    - **Financial transactions** (e.g., stock trading)
    - **Gaming** (e.g., multiplayer online games)
    - **Streaming services** (e.g., live video, audio)
- 3) **Improves user experience**
- 4) **Increases efficiency**

## How to Achieve Low Latency?

- 1) **Caching**
  - Store frequently accessed data in memory instead of retrieving it from a slow database.
  - Examples:
    - **Redis**
    - **Memcached**
- 2) **Optimized Database Queries**
  - Use indexes for faster lookups.
  - Implement connection pooling.
  - Avoid expensive joins and unnecessary queries.
- 3) **Load Balancer**
  - Distribute traffic across multiple servers to prevent overload on a single server.
- 4) **CDN (Content Delivery Network)**
  - Serve static assets (images, videos, etc.) from the closest location to the user.

## High Availability

- **Availability**
  - The ability to remain accessible.
- **High Availability (HA)**
  - A system is reliable and accessible with **minimal downtime**.

## Why is High Availability Important?

- 1) Users expect **24/7 uptime**,
  - even on high-traffic days (e.g., Black Friday).
- 2) Ensures **business continuity**
  - prevents revenue loss
- 3) Reduces **downtime costs**
  - losses due to downtime

- 4) Builds **customer trust**.

### **How to Achieve High Availability?**

- 1) **Failover Mechanisms**
  - Run multiple instances of a service.
  - If one instance fails, another should take over.
- 2) **Load Balancer**
  - Distributes traffic to multiple servers to prevent overloading a single server.
- 3) **Database Replication**
  - Have multiple copies of the database for failover protection.
- 4) **Auto-Scaling**
  - Dynamically increase or decrease the number of instances based on traffic.
- 5) **Graceful Degradation**
  - If part of the system fails, reduce functionality instead of going completely offline.